

Introduction to MongoDB

MongoDB: Introduction

- The leader in the NoSQL Document-based databases
- Full of features, beyond NoSQL
 - High performance
 - High availability
 - Native scalability
 - High flexibility
 - Open source

Terminology – Approximate mapping

Relational database	MongoDB
Table	Collection
Record	Document
Column	Field

MongoDB: Document Data Design

- High-level, business-ready representation of the data
- Records are stored into Documents
 - field-value pairs
 - similar to JSON objects
 - may be nested

```
{
  _id: <ObjectID1>,
  username: "123xyz",
  contact: {
    phone: 1234567890,
    email: "xyz@email.com",
  },
  access: {
    level: 5,
    group: "dev",
  }
}
```

Embedded Sub-Document

Embedded Sub-Document

MongoDB: Document Data Design

- High-level, business-ready representation of the data
- Flexible and rich syntax, adapting to most use cases
- Mapping into developer-language objects
 - year, month, day, timestamp,
 - lists, sub-documents, etc.

MongoDB: Main features

➤ Rich query language

- Documents can be created, read, updated and deleted.
- The **SQL language** is **not supported**
- APIs available for many programming languages
 - JavaScript, PHP, Python, Java, C#, ..

MongoDB: query language

➤ Most of the operations available in SQL language can be expressend in MongoDB language

MySQL clause	MongoDB operator
SELECT	<code>find()</code>

SELECT * FROM people	<code>db.people.find()</code>
--------------------------------	-------------------------------

MongoDB: Read data from documents

➤ Select documents

- `db.<collection name>.find({<conditions>}, {<fields of interest>});`

➤ E.g.,

`db.people.find();`

- Returns all documents contained in the people collection

MongoDB: Read data from documents

➤ Select documents

- `db.<collection name>.find({<conditions>}, {<fields of interest>});`

➤ Select the documents satisfying the specified conditions and specifically only the fields specified in fields of interest

- `<conditions>` are optional
 - conditions take a document with the form:
`{field1 : <value>, field2 : <value> ... }`
 - Conditions may specify a value or a regular expression

MongoDB: Read data from documents

➤ Select documents

- `db.<collection name>.find({<conditions>}, {<fields of interest>});`

➤ Select the documents satisfying the specified conditions and specifically only the fields specified in fields of interest

- `<fields of interest>` are optional
 - projections take a document with the form:
`{field1 : <value>, field2 : <value> ... }`
 - 1/true to include the field, 0/false to exclude the field

MongoDB: Read data from documents

➤ E.g.,

```
db.people.find().pretty();
```

- No conditions and no fields of interest
 - Returns all documents contained in the people collection
 - `pretty()` displays the results in an easy-to-read format

```
db.people.find({age:55})
```

- One condition on the value of age
 - Returns all documents having *age* equal to 55

MongoDB: Read data from documents

```
db.people.find({ }, { user_id: 1, status: 1 })
```

➤ No conditions, but returns a specific set of fields of interest

- Returns only **user_id** and **status** of all documents contained in the people collection
- Default of fields is false, except for _id

```
db.people.find({ status: "A", age: 55 })
```

➤ status = "A" and age = 55

- Returns all documents having **status = "A"** and **age = 55**

MongoDB: find() operator

MySQL clause	MongoDB operator
SELECT	find()

<pre>SELECT id, user_id, status FROM people</pre>	<pre>db.people.find({ }, { user_id: 1, status: 1 })</pre>
---	---

MongoDB: find() operator

MySQL clause	MongoDB operator
SELECT	find()

<pre>SELECT id, user_id, status FROM people</pre>	<pre>db.people.find({ }, { user_id: 1, status: 1 })</pre>
---	---

Where Condition

Select fields

MongoDB: find() operator

MySQL clause	MongoDB operator
SELECT	find()
WHERE	find({<WHERE CONDITIONS>})

<pre>SELECT * FROM people WHERE status = "A"</pre>	<pre>db.people.find({ status: "A" })</pre>
--	--



Where Condition

MongoDB: find() operator

MySQL clause	MongoDB operator
SELECT	find()
WHERE	find({<WHERE CONDITIONS>})

Where Condition

<pre>SELECT user_id, status FROM people WHERE status = "A"</pre>	<pre>db.people.find({ status: "A" }, { user_id: 1, status: 1, _id: 0 })</pre>
--	---

By default, the `_id` field is shown.
To remove it from visualization use: `_id: 0`

Selection fields

MongoDB: find() operator

MySQL clause	MongoDB operator
SELECT	find()
WHERE	find({<WHERE CONDITIONS>})

	<pre>db.people.find({"address.city": "Rome" })</pre>
--	--

```
{ _id: "A",  
  address: {  
    street: "Via Torino",  
    number: "123/B",  
    city: "Rome",  
    code: "00184"  
  }  
}
```

nested document



MongoDB: Read data from documents

```
db.people.find({ age: { $gt: 25, $lte: 50 } })
```

➤ Age greater than 25 and less than or equal to 50

- Returns all documents having **age > 25 and age <= 50**

```
db.people.find({$or:[{status: "A"},{age: 55}]})
```

➤ Status = "A" or age = 55

- Returns all documents having **status="A" or age=55**

```
db.people.find({ status: {$in:["A", "B"]}})
```

➤ Status = "A" or status = B

- Returns all documents where the **status** field value is **either "A" or "B"**

MongoDB: Read data from documents

➤ Select a single document

- `db.<collection name>.findOne({<conditions>}, {<fields of interest>});`

➤ Select one document that satisfies the specified query criteria.

- If multiple documents satisfy the query, it returns the first one according to the natural order which reflects the order of documents on the disk.

MongoDB: (no) joins

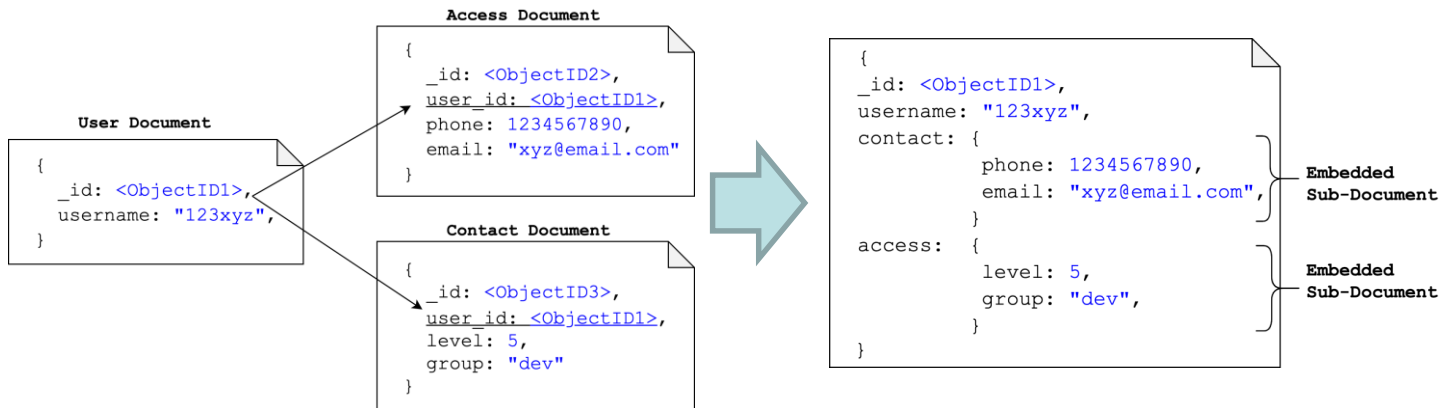
- There are other operators for selecting data from MongoDB collections
- However, no join operator exists (but `$lookup`)
 - You must write a program that
 - Selects the documents of the first collection you are interested in
 - Iterates over the documents returned by the first step, by using the loop statement provided by the programming language you are using
 - Executes one query for each of them to retrieve the corresponding document(s) in the other collection

<https://docs.mongodb.com/manual/reference/operator/aggregation/lookup>

MongoDB: (no) joins

➤ (no) joins

- Relations among documents/records are provided by
 - Object(ID) reference, with **no native join**
 - **DBRef**, across collections and databases



MongoDB: comparison operators

- In SQL language, comparison operators are essential to express conditions on data.
- In Mongo query language they are available with a different syntax.

MySQL	MongoDB	Description
>	\$gt	greater than
>=	\$gte	greater equal then
<	\$lt	less than
<=	\$lte	less equal then
=	\$eq	equal to
!=	\$neq	not equal to

MongoDB: Comparison query operators

Name	Description
<code>\$eq</code> or <code>:</code>	Matches values that are equal to a specified value
<code>\$gt</code>	Matches values that are greater than a specified value
<code>\$gte</code>	Matches values that are greater than or equal to a specified value
<code>\$in</code>	Matches any of the values specified in an array
<code>\$lt</code>	Matches values that are less than a specified value
<code>\$lte</code>	Matches values that are less than or equal to a specified value
<code>\$ne</code>	Matches all values that are not equal to a specified value
<code>\$nin</code>	Matches none of the values specified in an array

MongoDB: comparison operators (>)

MySQL	MongoDB	Description
>	\$gt	greater than

<pre>SELECT * FROM people WHERE age > 25</pre>	<pre>db.people.find({ age: { \$gt: 25 } })</pre>
---	--

MongoDB: comparison operators (>=)

MySQL	MongoDB	Description
>	\$gt	greater than
>=	\$gte	greater equal then

```
SELECT *  
FROM people  
WHERE age >= 25
```

```
db.people.find(  
  { age: { $gte: 25 } }  
)
```

MongoDB: comparison operators (<)

MySQL	MongoDB	Description
>	\$gt	greater than
>=	\$gte	greater equal then
<	\$lt	less than

```
SELECT *  
FROM people  
WHERE age < 25
```

```
db.people.find(  
  { age: { $lt: 25 } }  
)
```

MongoDB: comparison operators (<=)

MySQL	MongoDB	Description
>	\$gt	greater than
>=	\$gte	greater equal then
<	\$lt	less than
<=	\$lte	less equal then

```
SELECT *  
FROM people  
WHERE age <= 25
```

```
db.people.find(  
  { age: { $lte: 25 } }  
)
```

MongoDB: comparison operators (=)

MySQL	MongoDB	Description
>	\$gt	greater than
>=	\$gte	greater equal then
<	\$lt	less than
<=	\$lte	less equal then
=	\$eq	equal to The \$eq expression is equivalent to { field: <value> }.

<pre>SELECT * FROM people WHERE age = 25</pre>	<pre>db.people.find({ age: { \$eq: 25 } })</pre>
---	---

MongoDB: comparison operators (!=)

MySQL	MongoDB	Description
>	\$gt	greater than
>=	\$gte	greater equal then
<	\$lt	less than
<=	\$lte	less equal then
=	\$eq	equal to
!=	\$neq	Not equal to

```
SELECT *  
FROM people  
WHERE age != 25
```

```
db.people.find(  
  { age: { $neq: 25 } } }  
)
```

MongoDB: conditional operators

- To specify multiple conditions, **conditional operators** are used
- MongoDB offers the same functionalities of MySQL with a different syntax.

MySQL	MongoDB	Description
AND	,	Both verified
OR	\$or	At least one verified

MongoDB: conditional operators (AND)

MySQL	MongoDB	Description
AND	,	Both verified

<pre>SELECT * FROM people WHERE status = "A" AND age = 50</pre>	<pre>db.people.find({ status: "A", age: 50 })</pre>
---	---

MongoDB: conditional operators (OR)

MySQL	MongoDB	Description
AND	,	Both verified
OR	<code>\$or</code>	At least one verified

<pre>SELECT * FROM people WHERE status = "A" OR age = 50</pre>	<pre>db.people.find({ <code>\$or</code>: [{ status: "A" } , { age: 50 }] })</pre>
--	---

MongoDB: Cursor

➤ `db.collection.find()` gives back a cursor. It can be used to iterate over the result or as input for next operations.

➤ E.g.,

- `cursor.sort()`
- `cursor.count()`
- `cursor.forEach()` //shell method
- `cursor.limit()`
- `cursor.max()`
- `cursor.min()`
- `cursor.pretty()`

➤ Cursor examples:

```
db.people.find({ status: "A" }).count()
```

- Select documents with status="A" and count them.

```
db.people.find({ status: "A" }).forEach(  
  function(myDoc) { print( "user: "+myDoc.name );  
  })
```

- forEach applies a JavaScript function to apply to each document from the cursor.
 - Select documents with status="A" and print the document name.

MongoDB: sorting data

➤ Sort is a cursor method

➤ Sort documents

- `sort({<list of field:value pairs>});`
- field specifies which field is used to sort the returned documents
- value = -1 descending order
- Value = 1 ascending order

➤ Multiple field: value pairs can be specified

- Documents are sort based on the first field
- In case of ties, the second specified field is considered

MongoDB: sorting data

➤ E.g.,

```
db.people.find({ status: "A" }).sort({ age: 1 })
```

- Select documents with status="A" and sort them in ascending order based on the age value
 - Returns all documents having status="A". The result is sorted in ascending age order

MongoDB: sorting data

➤ Sorting data with respect to a given field in MongoDB: `sort()` operator

MySQL clause	MongoDB operator
ORDER BY	<code>sort()</code>

<pre>SELECT * FROM people WHERE status = "A" ORDER BY user_id ASC</pre>	<pre>db.people.find({ status: "A" }) .sort({ user_id: 1 })</pre>
---	--

MongoDB: sorting data

➤ Sorting data with respect to a given field in MongoDB: `sort()` operator

MySQL clause	MongoDB operator
ORDER BY	<code>sort()</code>

<pre>SELECT * FROM people WHERE status = "A" ORDER BY user_id ASC</pre>	<pre>db.people.find({ status: "A" }).sort({ user_id: 1 })</pre>
<pre>SELECT * FROM people WHERE status = "A" ORDER BY user_id DESC</pre>	<pre>db.people.find({ status: "A" }).sort({ user_id: -1 })</pre>

MongoDB: counting

MySQL clause	MongoDB operator
COUNT	<code>count()</code> or <code>find().count()</code>

<pre>SELECT COUNT(*) FROM people</pre>	<pre>db.people.count() or db.people.find().count()</pre>
--	--

MongoDB: counting

MySQL clause	MongoDB operator
COUNT	<code>count()</code> or <code>find().count()</code>

➤ Similar to the `find()` operator, `count()` can embed conditional statements.

<pre>SELECT COUNT(*) FROM people WHERE age > 30</pre>	<pre>db.people.count({ age: { \$gt: 30 } })</pre>
---	--

Introduction to data aggregation

Aggregation in MongoDB

- Aggregation operations process data records and return computed results.
- Documents enter a multi-stage pipeline that transforms the documents into an aggregated result.

MongoDB: Aggregation Framework

SQL	MongoDB
WHERE	\$match
GROUP BY	\$group
HAVING	\$match
SELECT	\$project
ORDER BY	\$sort
<u>//LIMIT</u>	<u>\$limit</u>
<u>SUM</u>	<u>\$sum</u>
<u>COUNT</u>	<u>\$sum</u>

MongoDB: Aggregation

➤ Aggregate functions can be applied to collections to group documents

```
db.collection.aggregate({<set of stages>})
```

- Common stages: `$match`, `$group` ..
- The aggregate function allows applying aggregating functions (e.g. `sum`, `average`, ..)
- It can be combined with an initial definition of groups based on the grouping fields

MongoDB: Aggregation

```
db.people.aggregate( [
  { $group: { _id: null,
              mytotal: { $sum: "$age" },
              mycount: { $sum: 1 }
            }
  ] )
```

- Considers all documents of people and
 - sum the values of their age
 - sum a set of ones (one for each document)
- The returned value is associated with a field called "mytotal" and a field "mycount"

MongoDB: Aggregation

```
db.people.aggregate( [
  { $group: { _id: null,
              myaverage: { $avg: "$age" },
              mytotal: { $sum: "$age" }
            }
  ] )
```

- Considers all documents of people and computes
 - sum of age
 - average of age

MongoDB: Aggregation

```
db.people.aggregate( [  
  { $match: {status: "A"} },  
  { $group: { _id: null,  
              count: { $sum: 1 }  
            }  
        }  
  ] )
```

Where conditions

- Counts the number of documents in people with status equal to "A"

MongoDB: Aggregation

```
db.people.aggregate( [
  { $group: { _id: "$status",
              count: { $sum: 1 }
            }
  ] )
```

- Creates one group of documents for each value of status and counts the number of documents per group
 - Returns one value for each group containing the value of the grouping field and an integer representing the number of documents


MongoDB: Aggregation

```
db.people.aggregate( [
  { $group: { _id: "$status",
              count: { $sum: 1 }
            }
  },
  { $match: { count: { $gte: 3 } } }
] )
```

- Creates one group of documents for each value of status and counts the number of documents per group. Returns only the groups with at least 3 documents

MongoDB: Aggregation

```
db.people.aggregate( [  
  { $group: { _id: "$status",  
              count: { $sum: 1 }  
            }  
  },  
  { $match: { count: { $gte: 3 } } }  
] )
```



The diagram illustrates the aggregation pipeline. A red box labeled "Having condition" points to the `{ $match: { count: { $gte: 3 } } }` stage, which filters the results to only include groups with a count of 3 or more documents.

- Creates one group of documents for each value of status and counts the number of documents per group. Returns only the groups with at least 3 documents

MongoDB: Aggregation Framework

SQL	MongoDB
WHERE	\$match
GROUP BY	\$group
HAVING	\$match
SELECT	\$project
ORDER BY	\$sort
LIMIT	\$limit
SUM	\$sum
COUNT	\$sum

Aggregation in MongoDB: Group By

MySQL clause	MongoDB operator
GROUP BY	aggregate(\$group)

```
SELECT status,  
       AVG(age) AS total  
FROM people  
GROUP BY status
```

```
db.orders.aggregate( [  
  {  
    $group: {  
      _id: "$status",  
      total: { $avg: "$age" }  
    }  
  }  
] )
```

Aggregation in MongoDB: Group By

MySQL clause	MongoDB operator
GROUP BY	aggregate(\$group)

```
SELECT status,  
       SUM(age) AS total  
FROM people  
GROUP BY status
```

```
db.orders.aggregate( [  
  {  
    $group: {  
      _id: "$status",  
      total: { $sum: "$age" }  
    }  
  }  
] )
```

Group field

Aggregation in MongoDB: Group By

MySQL clause	MongoDB operator
GROUP BY	aggregate(\$group)

```
SELECT status,  
       SUM(age) AS total  
FROM people  
GROUP BY status
```

```
db.orders.aggregate( [  
  {  
    $group: {  
      _id: "$status",  
      total: { $sum: "$age" }  
    }  
  }  
]
```

Group field

Aggregation function

Aggregation in MongoDB: Group By

MySQL clause	MongoDB operator
HAVING	aggregate(\$group, \$match)

```
SELECT status,  
       SUM(age) AS total  
FROM people  
GROUP BY status  
HAVING total > 1000
```

```
db.orders.aggregate( [  
  {  
    $group: {  
      _id: "$status",  
      total: { $sum: "$age" }  
    }  
  },  
  { $match: { total: { $gt: 1000 } } }  
)
```

Aggregation in MongoDB: Group By

MySQL clause	MongoDB operator
HAVING	aggregate(\$group, \$match)

```
SELECT status,  
       SUM(age) AS total  
FROM people  
GROUP BY status  
HAVING total > 1000
```

```
db.orders.aggregate( [
```

```
  {  
    $group: {  
      _id: "$status",  
      total: { $sum: "$age" }  
    }  
  },
```

```
  { $match: { total: { $gt: 1000 } } }  
] )
```

Group stage: Specify the aggregation field and the aggregation function

Aggregation in MongoDB: Group By

MySQL clause	MongoDB operator
HAVING	aggregate(\$group, \$match)

```
SELECT status,  
       SUM(age) AS total  
FROM people  
GROUP BY status  
HAVING total > 1000
```

```
db.orders.aggregate( [  
  {  
    $group: {  
      _id: "$status",  
      total: { $sum: "$age" }  
    }  
  },  
  { $match: { total: { $gt: 1000 } }  
}] )
```

Group stage: Specify the aggregation field and the aggregation function

Match Stage: specify the condition as in HAVING

Aggregation in MongoDB

Collection
↓
`db.orders.aggregate(
 $match phase → { $match: { status: "A" } },
 $group phase → { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }
)`

{ cust_id: "A123", amount: 500, status: "A" }
{ cust_id: "A123", amount: 250, status: "A" }
{ cust_id: "B212", amount: 200, status: "A" }
{ cust_id: "A123", amount: 300, status: "D" }

orders

→
\$match

{ cust_id: "A123", amount: 500, status: "A" }
{ cust_id: "A123", amount: 250, status: "A" }
{ cust_id: "B212", amount: 200, status: "A" }

→
\$group

Results	
{	<code>_id: "A123",</code> total: 750
}	
<hr/>	
{	<code>_id: "B212",</code> total: 200
}	

GUI for Mongo DB

MongoDB Compass

- Visually explore data.
- Available on Linux, Mac, or Windows.
- MongoDB Compass analyzes documents and displays rich structures within collections.
- Visualize, understand, and work with your geospatial data.



MongoDB Compass

The screenshot shows the 'Connect to Host' window in MongoDB Compass. The window has a title bar 'MongoDB Compass - Connect'. On the left is a sidebar with options: 'CREATE FREE ATLAS CLUSTER' (with a link to 'Learn more'), 'New Connection', 'Favorites', and 'RECENTS'. The 'RECENTS' list shows several connections to 'bigdatadb.polito.it:27017' with timestamps. The main area contains the following fields:

- Hostname:** bigdatadb.polito.it
- Port:** 27017
- SRV Record:** A toggle switch that is currently turned off.
- Authentication:** A dropdown menu set to 'Username / Password'.
- Username:** Gestionali
- Password:** A field with masked characters (dots).
- Authentication Database:** dbdmg
- Replica Set Name:** An empty text field.
- Read Preference:** A dropdown menu set to 'Primary'.
- SSL:** A dropdown menu set to 'Unvalidated (insecure)'.
- SSH Tunnel:** A dropdown menu set to 'None'.
- Favorite Name:** A field containing 'e.g. Shared Dev, QA Box, PRODUCTION'.

➤ Connect to local or remote instances of MongoDB.

MongoDB Compass

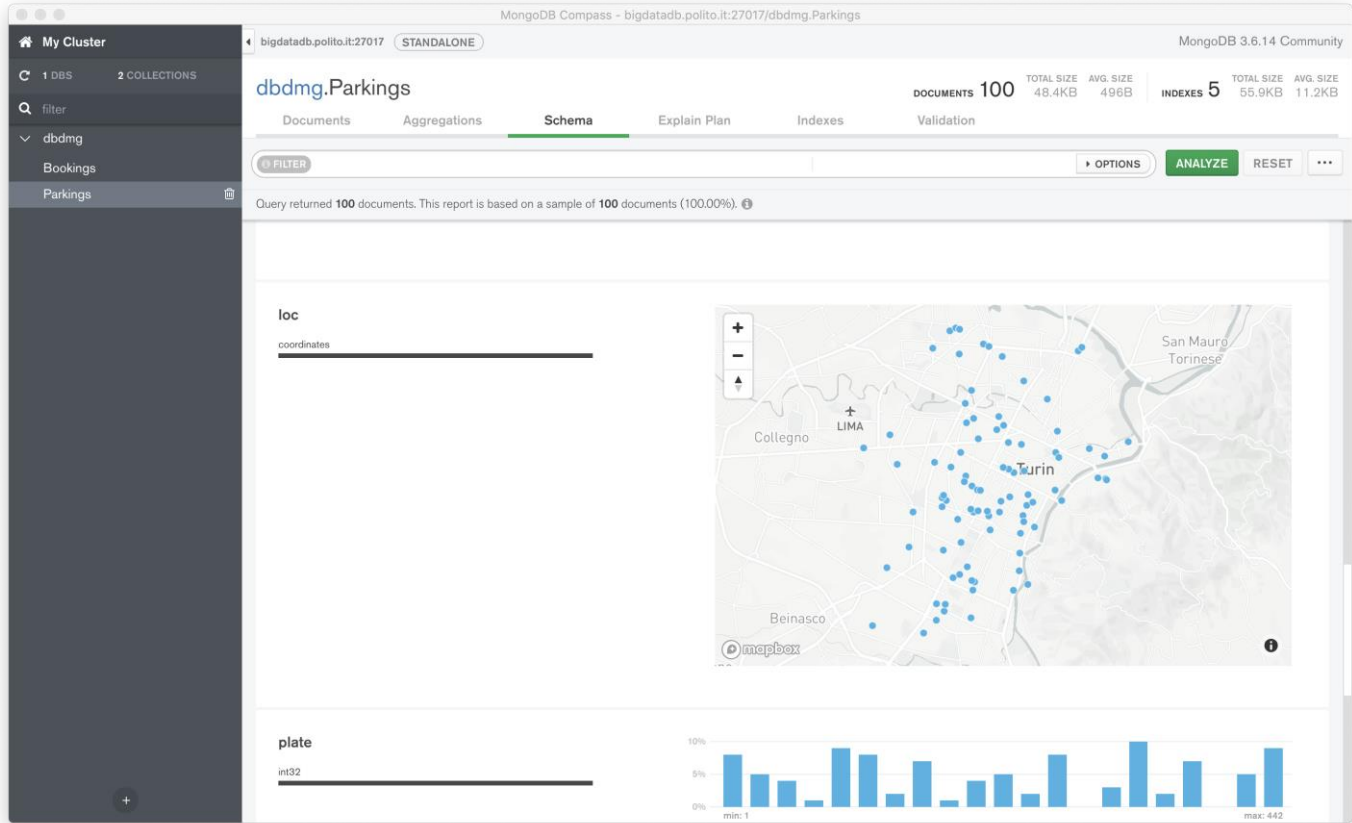
The screenshot displays the MongoDB Compass interface for a cluster named 'bigdatadb.polito.it:27017'. The selected database is 'dbdmg' and the collection is 'Parkings'. The interface shows a list of documents in a table format. The table has two columns: '_id' and 'ObjectId'. The documents are listed with their unique identifiers and corresponding ObjectId values.

_id	ObjectId
1	59bef0cd2ad8532c2a60093d
2	59bef0cd2ad8532c2a60093d
3	59bef1952ad8532c2a60093d
4	59bef1c62ad8532c2a60093d
5	59bef25c2ad8532c2a60093d
6	59bef25c2ad8532c2a60093d
7	59bef25c2ad8532c2a60093d
8	59bef25c2ad8532c2a60093d
9	59bef2bd2ad8532c2a60093d
10	59bef2bd2ad8532c2a60093d
11	59bef31d2ad8532c2a60093d
12	59bef34e2ad8532c2a60093d
13	59bef34e2ad8532c2a60093d
14	59bef37e2ad8532c2a60093d
15	59bef37e2ad8532c2a60093d
16	59bef37e2ad8532c2a60093d
17	59bef3e22ad8532c2a60093d
18	59bef3e22ad8532c2a60093d
19	59bef3e22ad8532c2a60093d
20	59bef4132ad8532c2a60093d

The interface also shows a sidebar with the cluster structure, including the database 'dbdmg' and its collections 'Bookings' and 'Parkings'. The main panel displays the 'Documents' tab, showing a list of documents with their unique identifiers and corresponding ObjectId values. The interface includes a search bar, a filter button, and a table view toggle. The status bar at the bottom indicates 'Displaying documents 1 - 20 of 100'.


➤ Get an overview of the data in list or table format.

MongoDB Compass



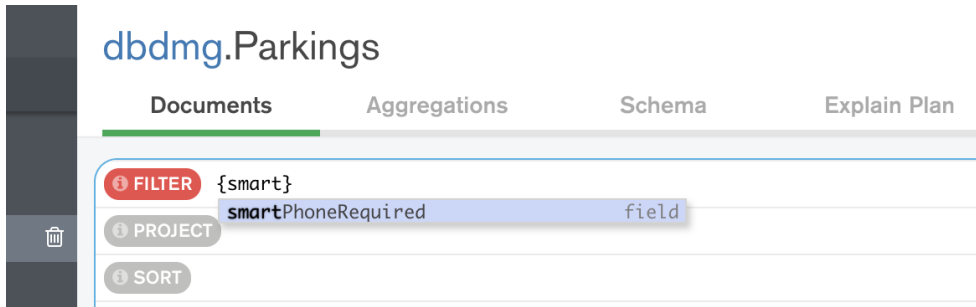
- Analyze the documents and their fields.
- Native support for geospatial coordinates.

MongoDB Compass

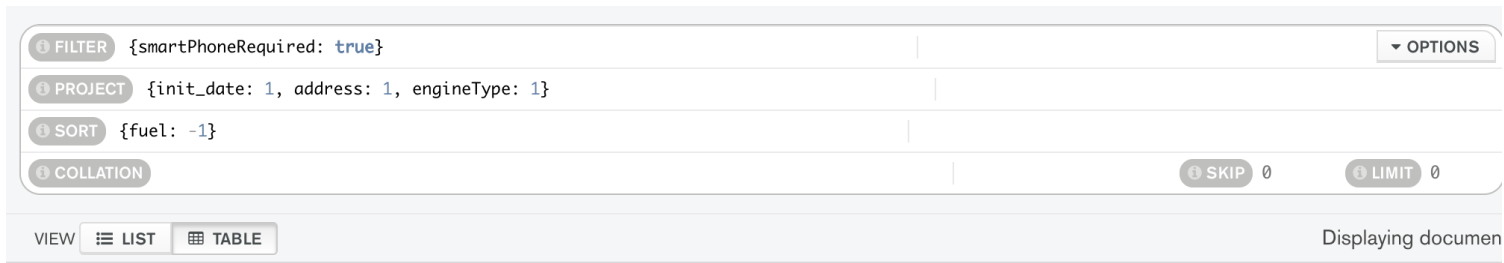


➤ Visually build the query conditioning on analyzed fields.

MongoDB Compass



➤ Autocomplete enabled by default.



➤ Construct the query step by step.

MongoDB Compass

The screenshot displays the MongoDB Compass interface for a cluster named 'bigdatadb.polito.it:27017'. The selected database is 'dbdmg' and the collection is 'Parkings'. The 'Explain Plan' tab is active, showing a query with a filter on 'interior' and a location-based query. The query performance summary indicates that 97 documents were returned, 0 index keys were examined, and 100 documents were examined. The actual query execution time was 0 ms, and the query was sorted in memory. A warning message states: 'No index available for this query.' The projection and sort stages are also shown, with the projection stage returning 97 documents and the sort stage returning 97 documents. The interface includes a sidebar with a cluster overview, a top navigation bar with tabs for Documents, Aggregations, Schema, Explain Plan, Indexes, and Validation, and a main content area with a query editor and performance summary.

Query Performance Summary

- Documents Returned: **97**
- Index Keys Examined: **0**
- Documents Examined: **100**
- Actual Query Execution Time (ms): **0**
- Sorted in Memory: **yes**
- ⚠️ No index available for this query.

PROJECTION

nReturned: **97** Execution Time: **0 ms**

Transform by:

```
{ "init_date": 1, "address": 1, "engineType": 1 }
```

SORT

nReturned: **97** Execution Time: **0 ms**

➤ Analyze query performance and get hints to speed it up.

MongoDB Compass

The screenshot shows the MongoDB Compass interface for a cluster named 'My Cluster'. The left sidebar shows the database 'dbdmg' and the collection 'Parkings'. The main panel is titled 'dbdmg.Parkings' and shows the 'Validation' tab. The 'Validation' tab displays the following information:

- Documents: 100
- Total Size: 48.4KB
- Avg. Size: 496B
- Indexes: 5
- Total Size: 55.9KB
- Avg. Size: 11.2KB

The 'Validation' tab includes a 'Validation Action' dropdown set to 'ERROR' and a 'Validation Level' dropdown set to 'STRICT'. Below these, a JSON schema is displayed:

```
1- {
2-   $jsonSchema: {
3-     required: ['exterior', 'interior', 'vendor', 'fuel'],
4-     properties: {
5-       vendor: {
6-         bsonType: "string",
7-         description: "must be a string"
8-       },
9-       fuel: {
10-        bsonType: "int",
11-        description: "must be an integer number"
12-      },
13-    }
14-  }
15- }
```

Below the schema, there are two sections:

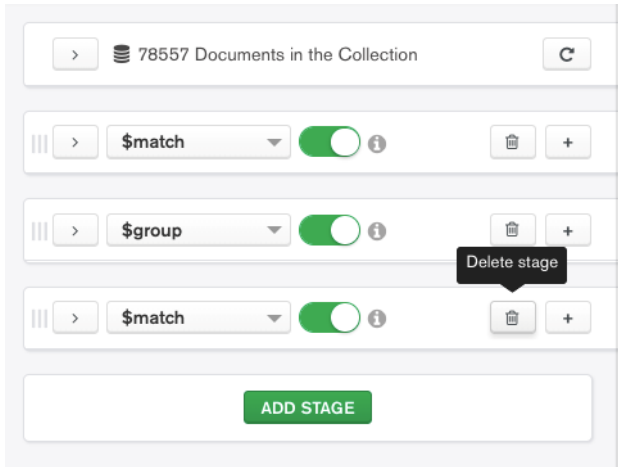
- Sample Document That Passed Validation**: A document with fields: `_id` (ObjectId), `plate` (442), `fuel` (37), `vendor` ("car2go"), `final_time` (1505685847), `loc` (Object), `init_time` (1505685697), `vin` ("VIN442"), and `smartPhoneReserved` (true).
- Sample Document That Failed Validation**: A section with the text "No Preview Documents".

At the bottom of the interface, there are two orange arrow icons pointing right, followed by the text "Specify constraints to validate data" and "Find inconsistent documents."


➤ Specify constraints to validate data

➤ Find inconsistent documents.

MongoDB Compass: Aggregation

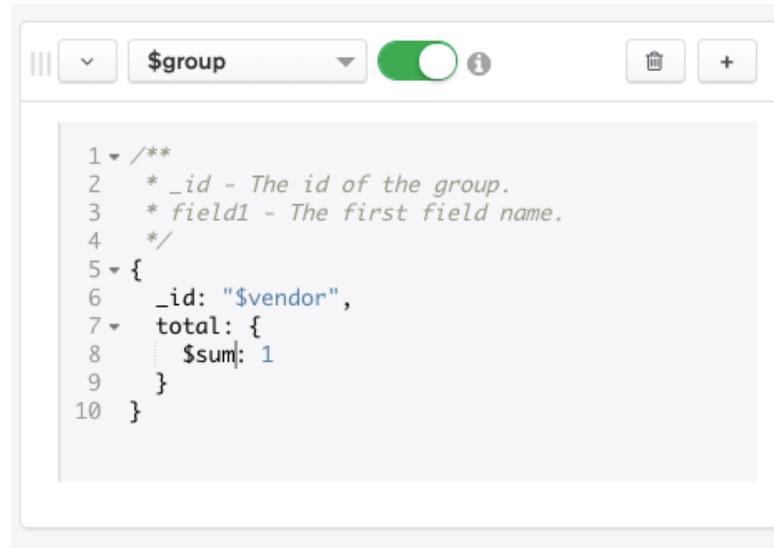


➤ Build a pipeline consisting of multiple aggregation stages.



➤ Define the filter and aggregation attributes for each operator.

MongoDB Compass: Aggregation stages



The screenshot shows the MongoDB Compass interface for editing an aggregation pipeline. The top bar includes a menu icon, a dropdown set to '\$group', a toggle switch, and an information icon. To the right are delete and add stage buttons. The main area contains a code editor with the following aggregation pipeline:

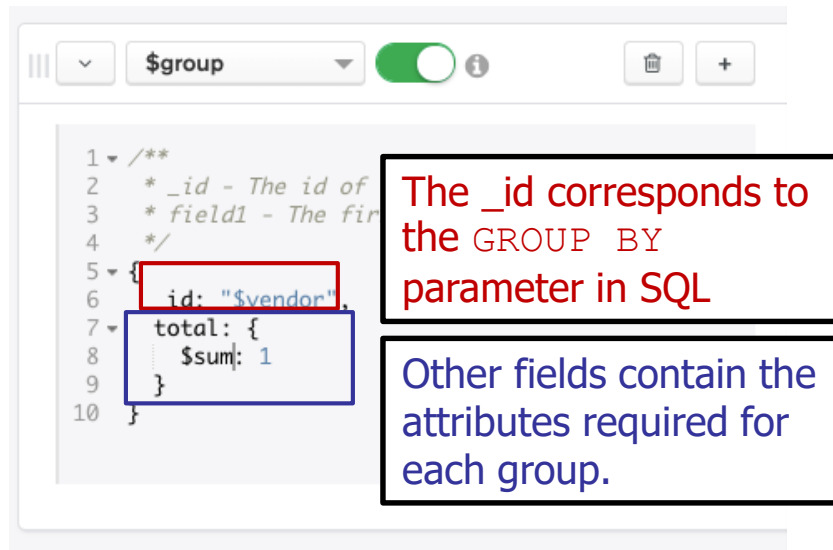
```
1 //  
2 * _id - The id of the group.  
3 * field1 - The first field name.  
4 */  
5 {  
6   _id: "$vendor",  
7   total: {  
8     $sum: 1  
9   }  
10 }
```

Output after \$group stage (Sample of 2 documents)

`_id: "car2go"`
`total: 48423`

`_id: "enjoy"`
`total: 30134`

MongoDB Compass: Aggregation stages



The screenshot shows the MongoDB Compass aggregation pipeline editor. The top bar includes a menu icon, a dropdown set to '\$group', a toggle switch, and icons for deleting and adding stages. The pipeline editor shows a single stage with the following code:

```
1 /**  
2  * _id - The id of  
3  * field1 - The fir  
4  */  
5 {  
6   id: "$vendor",  
7   total: {  
8     $sum: 1  
9   }  
10 }
```

Two callout boxes provide additional context:

- The `_id` corresponds to the GROUP BY parameter in SQL** (Red text, black border)
- Other fields contain the attributes required for each group.** (Blue text, black border)



The screenshot shows the output of the aggregation pipeline. The title is "Output after \$group stage (Sample of 2 documents)". The output is displayed in two separate boxes, each representing a group:

- `_id: "car2go"`
`total: 48423`
- `_id: "enjoy"`
`total: 30134`

One group for each "vendor".

MongoDB Compass: Pipelines

The screenshot shows the MongoDB Compass interface with a pipeline of two stages. The first stage is a `$group` operation, and the second stage is a `$match` operation. An arrow points from the `$group` stage to the `$match` stage, indicating the flow of the pipeline.

1st stage: grouping by vendor

```
1 /**
2  * _id - The id of the group.
3  * field1 - The first field name.
4  */
5 {
6   _id: "$vendor",
7   total: { $sum: 1 },
8   avg_fuel: { $avg: "$fuel" }
9 }
10
```

Output after \$group stage (Sample of 2 documents)

_id	total	avg_fuel
car2go	48423	64.88284492906264
enjoy	30134	61.03381562354815

2nd stage: condition over fields created in the previous stage (avg_fuel, total).

```
1 /**
2  * query - The query in MQL.
3  */
4 {
5   avg_fuel: { $gt: 63 },
6   total: { $gt: 35000 }
7 }
```

Output after \$match stage (Sample of 1 document)

_id	total	avg_fuel
car2go	48423	64.88284492906264

Indexing

MongoDB: Indexes

- Indexes are data structures that store a small portion of the collection's data set in a form easy to traverse.
- They store ordered values of a specific field, or set of fields, in order to efficiently support equality matches, range-based queries and sorting operations.

MongoDB: Indexes

- MongoDB provides different data-type indexes
- Single field indexes
 - Compound field indexes
 - Multikey indexes
 - Geospatial indexes
 - Text indexes
 - Hashed indexes

MongoDB: Create new indexes

➤ Creating an index

```
db.collection.createIndex(<index keys>, <options>)
```

- Before v. 3.0 use `db.collection.ensureIndex()`

➤ Options include: `name`, `unique` (whether to accept or not insertion of documents with duplicate index keys), `background`, `dropDups`, ..

MongoDB: Indexes

➤ Single field indexes

- Support user-defined ascending/descending indexes on a single field of a document

➤ E.g.,

- `db.orders.createIndex({orderDate: 1})`

➤ Compound field indexes

- Support user-defined indexes on a set of fields

➤ E.g.,

- `db.orders.createIndex({orderDate: 1,
zipcode: -1})`

MongoDB: Indexes

➤ MongoDB supports efficient queries of geospatial data

➤ Geospatial data are stored as:

- GeoJSON objects: embedded document { <type>, <coordinate> }
 - E.g., location: {type: "Point", coordinates: [-73.856, 40.848]}
- Legacy coordinate pairs: array or embedded document
 - point: [-73.856, 40.848]

MongoDB: Indexes

➤ Geospatial indexes

- Two type of geospatial indexes are provided: `2d` and `2dsphere`

➤ A `2dsphere` index supports queries that calculate geometries on an earth-like sphere

➤ Use a `2d` index for data stored as points on a two-dimensional plane.

➤ E.g.,

- `db.places.createIndex({location: "2dsphere"})`

➤ Geospatial query operators

- `$geoIntersects`, `$geoWithin`, `$near`, `$nearSphere`

MongoDB: Indexes

➤ \$near syntax:

```
{
  <location field>: {
    $near: {
      $geometry: {
        type: "Point" ,
        coordinates: [ <longitude> , <latitude> ]
      },
      $maxDistance: <distance in meters>,
      $minDistance: <distance in meters>
    }
  }
}
```

MongoDB: Indexes

➤ E.g.,

- `db.places.createIndex({location: "2dsphere"})`

➤ Geospatial query operators

- `$geoIntersects`, `$geoWithin`, `$near`, `$nearSphere`

➤ Geospatial aggregation stage

- `$near`

MongoDB: Indexes

➤ E.g.,

- ```
db.places.find({location:
 {$near:
 {$geometry: {
 type: "Point",
 coordinates: [-73.96, 40.78] },
 $maxDistance: 5000}
 })
```
- Find all the places within 5000 meters from the specified GeoJSON point, sorted in order from nearest to furthest

# MongoDB: Indexes

## ➤ Text indexes

- Support efficient searching for string content in a collection
- Text indexes store only *root words* (no language-specific *stop words* or *stem*)

## ➤ E.g.,

```
db.reviews.createIndex({comment: "text"})
```

- Wildcard (\$\*\*) allows MongoDB to index every field that contains string data
- E.g.,

```
db.reviews.createIndex({"$**": "text"})
```